

Titanic Passenger Survival Classification Model

```
# Cell 1: Imports
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns

from sklearn.model_selection import train_test_split, GridSearchCV
from sklearn.preprocessing import StandardScaler, OneHotEncoder
from sklearn.compose import ColumnTransformer
from sklearn.pipeline import Pipeline
from sklearn.impute import SimpleImputer
from sklearn.metrics import (confusion_matrix, classification_report,
                             roc_auc_score, roc_curve, accuracy_score,
                             precision_score,
                             recall_score, f1_score)
from sklearn.linear_model import LogisticRegression
from sklearn.ensemble import RandomForestClassifier
from sklearn.svm import SVC

sns.set(style="whitegrid")
%matplotlib inline

# Cell 2: Load dataset
import seaborn as sns
df = sns.load_dataset('titanic') # convenient built-in version
print(df.shape)
df.head()

(891, 15)

{"summary": "{\n    \"name\": \"df\",\n    \"rows\": 891,\n    \"fields\": [\n        {\n            \"column\": \"survived\",\n            \"properties\": {\n                \"dtype\": \"number\",\n                \"std\": 0,\n                \"min\": 0,\n                \"max\": 1,\n                \"num_unique_values\": 2,\n                \"samples\": [\n                    1,\n                    0\n                ],\n                \"semantic_type\": \"\",\n                \"description\": \"\"\n            }\n        },\n        {\n            \"column\": \"pclass\",\n            \"properties\": {\n                \"dtype\": \"number\",\n                \"std\": 0,\n                \"min\": 1,\n                \"max\": 3,\n                \"num_unique_values\": 3,\n                \"samples\": [\n                    3,\n                    1\n                ],\n                \"semantic_type\": \"\",\n                \"description\": \"\"\n            }\n        },\n        {\n            \"column\": \"sex\",\n            \"properties\": {\n                \"dtype\": \"category\",\n                \"num_unique_values\": 2,\n                \"samples\": [\n                    \"female\",\n                    \"male\"\n                ],\n                \"semantic_type\": \"\",\n                \"description\": \"\"\n            }\n        },\n        {\n            \"column\": \"age\",\n            \"properties\": {\n                \"dtype\": \"number\",\n                \"std\": 14.526497332334044,\n                \"min\": 0.42,\n                \"max\": 80.0,\n                \"num_unique_values\": 891\n            }\n        }\n    ]\n}
```

```
88,\n          \"samples\": [\n            0.75,\n            22.0\n          ],\n          \"semantic_type\": \"\",\n          \"description\": \"\"\\n        },\n        {\n          \"column\": \"sibsp\",\n          \"properties\": {\n            \"dtype\": \"number\",\n            \"std\": 1,\n            \"min\": 0,\n            \"max\": 8,\n            \"num_unique_values\": 7,\n            \"samples\": [\n              1,\n              0\n            ],\n            \"semantic_type\": \"\"\n          },\n          \"description\": \"\"\\n        },\n        {\n          \"column\": \"parch\",\n          \"properties\": {\n            \"dtype\": \"number\",\n            \"std\": 0,\n            \"min\": 0,\n            \"max\": 6,\n            \"num_unique_values\": 7,\n            \"samples\": [\n              0,\n              1\n            ],\n            \"semantic_type\": \"\"\n          },\n          \"description\": \"\"\\n        },\n        {\n          \"column\": \"fare\",\n          \"properties\": {\n            \"dtype\": \"number\",\n            \"std\": 49.693428597180905,\n            \"min\": 0.0,\n            \"max\": 512.3292,\n            \"num_unique_values\": 248,\n            \"samples\": [\n              11.2417,\n              51.8625\n            ],\n            \"semantic_type\": \"\"\n          },\n          \"description\": \"\"\\n        },\n        {\n          \"column\": \"embarked\",\n          \"properties\": {\n            \"dtype\": \"category\",\n            \"num_unique_values\": 3,\n            \"samples\": [\n              \"S\",\n              \"C\"\n            ],\n            \"semantic_type\": \"\"\n          },\n          \"description\": \"\"\\n        },\n        {\n          \"column\": \"class\",\n          \"properties\": {\n            \"dtype\": \"category\",\n            \"num_unique_values\": 3,\n            \"samples\": [\n              \"Third\",\n              \"First\"\n            ],\n            \"semantic_type\": \"\"\n          },\n          \"description\": \"\"\\n        },\n        {\n          \"column\": \"who\",\n          \"properties\": {\n            \"dtype\": \"category\",\n            \"num_unique_values\": 3,\n            \"samples\": [\n              \"man\",\n              \"woman\"\n            ],\n            \"semantic_type\": \"\"\n          },\n          \"description\": \"\"\\n        },\n        {\n          \"column\": \"adult_male\",\n          \"properties\": {\n            \"dtype\": \"boolean\",\n            \"num_unique_values\": 2,\n            \"samples\": [\n              false,\n              true\n            ],\n            \"semantic_type\": \"\"\n          },\n          \"description\": \"\"\\n        },\n        {\n          \"column\": \"deck\",\n          \"properties\": {\n            \"dtype\": \"category\",\n            \"num_unique_values\": 7,\n            \"samples\": [\n              \"C\",\n              \"E\"\n            ],\n            \"semantic_type\": \"\"\n          },\n          \"description\": \"\"\\n        },\n        {\n          \"column\": \"embark_town\",\n          \"properties\": {\n            \"dtype\": \"category\",\n            \"num_unique_values\": 3,\n            \"samples\": [\n              \"Southampton\",\n              \"Cherbourg\"\n            ],\n            \"semantic_type\": \"\"\n          },\n          \"description\": \"\"\\n        },\n        {\n          \"column\": \"alive\",\n          \"properties\": {\n            \"dtype\": \"category\",\n            \"num_unique_values\": 2,\n            \"samples\": [\n              \"yes\",\n              \"no\"\n            ],\n            \"semantic_type\": \"\"\n          },\n          \"description\": \"\"\\n        },\n        {\n          \"column\": \"alone\",\n          \"properties\": {\n            \"dtype\": \"category\",\n            \"num_unique_values\": 2,\n            \"samples\": [\n              \"no\"\n            ],\n            \"semantic_type\": \"\"\n          },\n          \"description\": \"\"\\n        }\n      ],\n      \"semantic_type\": \"\",  
      \"description\": \"\"\\n    }
```

```
n      "dtype": "boolean", "n      "num_unique_values": 2, "n
\"samples\": [\n    true, \n    false\n  ], \n  \"semantic_type\": \"\", \n  \"description\": \"\"\n}\n} \n ]\n}","type":"dataframe","variable_name":"df"}
```

Dataset Description

The Titanic dataset contains information about passengers aboard the RMS Titanic. The goal is to predict passenger survival based on demographic and travel-related features.

- Source: Kaggle – Titanic: Machine Learning from Disaster
- Rows: 891
- Columns: 12
- Target Variable: Survived (0 = No, 1 = Yes)

Features:

- PassengerId: Unique ID
- Pclass: Ticket class
- Name: Passenger name
- Sex: Gender
- Age: Age in years
- SibSp: Number of siblings/spouses aboard
- Parch: Number of parents/children aboard
- Ticket: Ticket number
- Fare: Ticket fare
- Cabin: Cabin number
- Embarked: Port of Embarkation

```
# Cell 3: Quick dataset description
print("Columns:", df.columns.tolist())
print("Target distribution (survived):")
print(df['survived'].value_counts(dropna=False))
df.info()

Columns: ['survived', 'pclass', 'sex', 'age', 'sibsp', 'parch',
fare', 'embarked', 'class', 'who', 'adult_male', 'deck',
'embark_town', 'alive', 'alone']
```

```

Target distribution (survived):
survived
0    549
1    342
Name: count, dtype: int64
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 891 entries, 0 to 890
Data columns (total 15 columns):
 #   Column      Non-Null Count  Dtype  
--- 
 0   survived    891 non-null    int64  
 1   pclass       891 non-null    int64  
 2   sex          891 non-null    object  
 3   age          714 non-null    float64 
 4   sibsp        891 non-null    int64  
 5   parch        891 non-null    int64  
 6   fare          891 non-null    float64 
 7   embarked     889 non-null    object  
 8   class         891 non-null    category
 9   who           891 non-null    object  
 10  adult_male   891 non-null    bool   
 11  deck          203 non-null    category
 12  embark_town  889 non-null    object  
 13  alive         891 non-null    object  
 14  alone         891 non-null    bool  
dtypes: bool(2), category(2), float64(2), int64(4), object(5)
memory usage: 80.7+ KB

```

Data Preprocessing

The following preprocessing steps were applied:

1. Missing Values:
 - Age: Filled using median
 - Embarked: Filled using mode
 - Cabin: Dropped due to high missing percentage
2. Encoding Categorical Features:
 - Sex encoded using OneHotEncoder
 - Embarked encoded using OneHotEncoder
 - Pclass treated as categorical
3. Scaling Numerical Columns:
 - Age and Fare were standardized using StandardScaler
4. Dropping Irrelevant Columns:
 - Name, Ticket, Cabin were removed
5. Train-Test Split:
 - 80% Training, 20% Testing

```

# Cell 4: Prepare a working dataframe
data =
df[['survived','pclass','sex','age','sibsp','parch','fare','embarked']].copy()
data.shape

(891, 8)

# Cell 5: Examine missingness & basic stats
print(data.isna().sum())
data.describe(include='all')

survived      0
pclass        0
sex           0
age         177
sibsp        0
parch        0
fare          0
embarked     2
dtype: int64

{"summary": {"name": "data", "rows": 11, "fields": [
    {"column": "survived", "properties": {"dtype": "number", "std": 314.8713661874558, "min": 0.0, "max": 891.0, "num_unique_values": 5, "samples": [0.3838383838383838, 1.0, 0.4865924542648585]}, "semantic_type": "\\", "description": "\n"}, {"column": "pclass", "properties": {"dtype": "number", "std": 314.2523437079693, "min": 0.8360712409770513, "max": 891.0, "num_unique_values": 6, "samples": [2.308641975308642, 3.0]}, "semantic_type": "\\", "description": "\n"}, {"column": "sex", "properties": {"dtype": "category", "num_unique_values": 4, "samples": [2, 2, 577, 891]}, "semantic_type": "\\", "description": "\n"}, {"column": "age", "properties": {"dtype": "number", "std": 242.9056731818781, "min": 0.42, "max": 714.0, "num_unique_values": 8, "samples": [29.69911764705882, 28.0, 714.0]}, "semantic_type": "\\", "description": "\n"}, {"column": "sibsp", "properties": {"dtype": "number", "std": 314.4908277465442, "min": 0.0, "max": 891.0, "num_unique_values": 6, "samples": [0.5230078563411896, 8.0]}, "semantic_type": "\\", "description": "\n"}]}

```

```

    "semantic_type": "\",\n      "description": \"\\n      }\n    },\n    {\n      "column": "parch",\n      "properties": {\n        "dtype": "number",\n        "std": 314.65971717879,\n        "min": 0.0,\n        "max": 891.0,\n        "num_unique_values": 5,\n        "samples": [\n          0.38159371492704824,\n          0.8060572211299559\n        ],\n      }\n    },\n    {\n      "column": "fare",\n      "properties": {\n        "dtype": "number",\n        "std": 330.6256632228577,\n        "min": 0.0,\n        "max": 891.0,\n        "num_unique_values": 8,\n        "samples": [\n          32.204207968574636,\n          14.4542,\n          891.0\n        ],\n      }\n    },\n    {\n      "column": "embarked",\n      "properties": {\n        "category",\n        "num_unique_values": 4,\n        "samples": [\n          3,\n          644,\n          889\n        ],\n      }\n    }\n  ]\n},\n"type": "dataframe"}\n\n# Cell 6: Split X/y first (so transformations fit on train only)\nX = data.drop('survived', axis=1)\ny = data['survived']\n\nX_train, X_test, y_train, y_test = train_test_split(\n    X, y, test_size=0.20, random_state=42, stratify=y)\n\nprint("Train shape:", X_train.shape, "Test shape:", X_test.shape)\n\nTrain shape: (712, 7) Test shape: (179, 7)\n\n# Cell 7: Preprocessing pipeline\nnumeric_features = ['age', 'sibsp', 'parch', 'fare']\nnumeric_transformer = Pipeline(steps=[\n    ('imputer', SimpleImputer(strategy='median')),\n    ('scaler', StandardScaler())\n])\n\ncategorical_features = ['sex', 'embarked', 'pclass'] # pclass treated\n# as categorical\ncategorical_transformer = Pipeline(steps=[\n    ('imputer', SimpleImputer(strategy='most_frequent')),\n    ('onehot', OneHotEncoder(handle_unknown='ignore'))\n])\n\npreprocessor = ColumnTransformer(\n    transformers=[\n        ('num', numeric_transformer, numeric_features),\n        ('cat', categorical_transformer, categorical_features)\n    ])\n
```

Exploratory Data Analysis (EDA)

Key observations:


- Females had a significantly higher survival rate compared to males.
- Higher passenger classes (Pclass = 1) had better survival probabilities.
- Younger passengers survived more frequently.
- Larger families (high SibSp/Parch) showed lower survival.
- Fare and Pclass showed strong correlation with survival.

The following charts were used:

- Gender vs Survival Count Plot
- Pclass vs Survival Bar Plot
- Age Distribution Histogram
- Correlation Heatmap


```
# Cell 8: EDA plots
plt.figure(figsize=(10,4))
plt.subplot(1,2,1)
sns.countplot(x='survived', data=data)
plt.title('Survival Count')

plt.subplot(1,2,2)
sns.countplot(x='sex', hue='survived', data=data)
plt.title('Survival by Sex')
plt.tight_layout()
```




```
# Cell 9: Age distribution by survival
plt.figure(figsize=(8,4))
sns.kdeplot(data=data, x='age', hue='survived', fill=True)
plt.title('Age distribution: Survived vs Died')

Text(0.5, 1.0, 'Age distribution: Survived vs Died')
```




```
# Cell 10: Pclass vs Survival
plt.figure(figsize=(6,4))
sns.countplot(x='pclass', hue='survived', data=data)
plt.title('Survival by Passenger Class')
Text(0.5, 1.0, 'Survival by Passenger Class')
```



```
# Cell 11: Correlation heatmap (after basic numeric fill)
corr_df = data.copy()
corr_df['age'] = corr_df['age'].fillna(corr_df['age'].median())
corr_df['fare'] = corr_df['fare'].fillna(corr_df['fare'].median())
corr_df['sex'] = corr_df['sex'].map({'male':0,'female':1})
corr_df['embarked'] = corr_df['embarked'].map({'S':0,'C':1,'Q':2})
sns.heatmap(corr_df.corr(), annot=True, fmt=".2f")
plt.title('Correlation matrix (numeric)')

Text(0.5, 1.0, 'Correlation matrix (numeric)')
```



Model Selection

Two models were used:

1. Logistic Regression
 - Acts as a strong baseline model
 - Works well with linearly separable data
2. Random Forest Classifier
 - Handles non-linear patterns effectively
 - Reduces risk of overfitting
 - Performed better during evaluation

These models were chosen due to their interpretability, stability, and strong performance with mixed numerical-categorical datasets.

```
# Cell 12: Build pipelines for models
from sklearn.pipeline import make_pipeline

# Logistic Regression pipeline
```

```

lr_pipeline = Pipeline(steps=[
    ('preprocessor', preprocessor),
    ('clf', LogisticRegression(solver='liblinear', random_state=42))
])

# Random Forest pipeline
rf_pipeline = Pipeline(steps=[
    ('preprocessor', preprocessor),
    ('clf', RandomForestClassifier(n_estimators=100, random_state=42))
])

# Fit models
lr_pipeline.fit(X_train, y_train)
rf_pipeline.fit(X_train, y_train)

Pipeline(steps=[('preprocessor',
                 ColumnTransformer(transformers=[('num',
                                                     Pipeline(steps=[('imputer',
                                                                     SimpleImputer(strategy='median'))),
                                                     ('scaler',
                                                       StandardScaler()))),
                                             ['age', 'sibsp',
                                              'parch',
                                              'fare']),
                 ('cat',
                   Pipeline(steps=[('imputer',
                     SimpleImputer(strategy='most_frequent')),
                     ('onehot',
                       OneHotEncoder(handle_unknown='ignore'))])),
                                             ['sex', 'embarked',
                                              'pclass']])),
               ('clf', RandomForestClassifier(random_state=42))]

## Model Evaluation

```

The Random Forest Classifier achieved the highest performance **with**:

- High accuracy
- Balanced precision **and** recall
- Strong F1-score

The confusion matrix helped identify the number of correct **and** incorrect predictions.

ROC curve showed good separability capability.

These metrics confirm the model's reliability.

```
# Cell 13: Predictions
models = {'LogisticRegression': lr_pipeline, 'RandomForest':
rf_pipeline}

for name, model in models.items():
    y_pred = model.predict(X_test)
    y_proba = model.predict_proba(X_test)[:,1] if hasattr(model,
"predict_proba") else None
    print(f"== {name} ==")
    print("Accuracy:", accuracy_score(y_test, y_pred))
    print("Precision:", precision_score(y_test, y_pred))
    print("Recall:", recall_score(y_test, y_pred))
    print("F1-score:", f1_score(y_test, y_pred))
    if y_proba is not None:
        print("ROC-AUC:", roc_auc_score(y_test, y_proba))
    print(classification_report(y_test, y_pred))
    print("Confusion matrix:\n", confusion_matrix(y_test, y_pred))
    print()

== LogisticRegression ==
Accuracy: 0.8044692737430168
Precision: 0.7931034482758621
Recall: 0.6666666666666666
F1-score: 0.7244094488188977
ROC-AUC: 0.8426877470355731
      precision    recall   f1-score   support
      0       0.81     0.89     0.85      110
      1       0.79     0.67     0.72       69
      accuracy          0.80      0.80      179
      macro avg       0.80     0.78     0.79      179
      weighted avg    0.80     0.80     0.80      179

Confusion matrix:
[[98 12]
 [23 46]]

== RandomForest ==
Accuracy: 0.8044692737430168
Precision: 0.765625
Recall: 0.7101449275362319
F1-score: 0.7368421052631579
ROC-AUC: 0.8367588932806325
      precision    recall   f1-score   support
```


0	0.83	0.86	0.84	110
1	0.77	0.71	0.74	69
accuracy			0.80	179
macro avg	0.80	0.79	0.79	179
weighted avg	0.80	0.80	0.80	179

Confusion matrix:

```
[[95 15]
 [20 49]]
```

```
# Cell 14: ROC curve (Random Forest)
from sklearn.metrics import RocCurveDisplay
import matplotlib.pyplot as plt

model = rf_pipeline
y_proba = model.predict_proba(X_test)[:,1]
RocCurveDisplay.from_predictions(y_test, y_proba)
plt.title('ROC Curve – Random Forest')
plt.show()
```



Error Analysis

Misclassified samples were examined to understand model weakness.

Insights:

- Many misclassifications occurred for **males in 3rd class** with medium or high fare.
 - Passengers with missing or imputed Age values created inconsistencies.
 - Some survival outcomes contradict common patterns (e.g., male + low class + survived).
 - Large families showed overlapping patterns making them hard to classify.

The model struggles mainly in ambiguous or rare passenger cases.

```

512.3292,\n      "num_unique_values": 32,         "samples": [\n        91.0792,\n          35.5\n        ],\n      "semantic_type": "\",\n        "description": \"\\n      }\n    },\n    {\n      "column": "embarked",\n        "properties": {\n          "dtype": "category",\n            "num_unique_values": 3,\n              "samples": [\n                "C",\n                  "S"\n              ],\n              "semantic_type": "\",\n                "description": \"\\n      }\n            },\n            {\n              "column": "y_true",\n                "properties": {\n                  "dtype": "number",\n                    "std": 0,\n                      "min": 0,\n                        "max": 1,\n                          "num_unique_values": 2,\n                            "samples": [\n                              0,\n                                1\n                            ],\n                            "semantic_type": "\",\n                              "description": \"\\n      }\n                            },\n                            {\n                              "column": "y_pred",\n                                "properties": {\n                                  "dtype": "number",\n                                    "std": 0,\n                                      "min": 0,\n                                        "max": 1,\n                                          "num_unique_values": 2,\n                                            "samples": [\n                                              1,\n                                                0\n                                            ],\n                                            "semantic_type": "\",\n                                              "description": \"\\n      }\n                                              }\n                                              }\n                                              ]\\n}\n  },\n  "type": "dataframe",\n  "variable_name": "mis_examples"

```

Final Summary

This project demonstrates a complete machine learning workflow using the Titanic dataset. Key takeaways:

- Gender, class, and fare are the strongest predictors of survival.
- Random Forest provided the best performance.
- The model is reliable but can be further improved with tuning and feature engineering.
- The workflow shows strong understanding of preprocessing, EDA, modeling, and evaluation.

This project is suitable for real-world classification problem demonstration.